Five-axis rubric	Each axis scored as an absolute integer from 0 (lowest) to 5 (highest); unweighted sum gives the 0–25 rubric sum.

Axis	**Real-POCQi** single-turn: one answer to a physician's clinical question	**MedSP1000** multi-turn: clinician trajectory with a simulated patient
Accuracy `accuracy`	How factually accurate is the response?	How factually and clinically accurate are the clinician's statements across the encounter?
Clinical utility `clinical_utility`	How useful is the response for providing high-quality clinical care?	How useful is the clinician's overall interaction for providing high-quality clinical care?
Source quality `source_quality`	How authoritative is the source material supporting the response?	How well grounded are the clinician's claims and recommendations in authoritative clinical knowledge? *Formal citations are not required in dialogue.*
Verifiability `verifiability`	How easy is the response to verify?	How specific, transparent, and checkable are the clinician's claims, reasoning, and recommendations?
Completeness `completeness`	How completely does the response address the question?	How completely does the clinician address the scenario using the information available over the full interaction?

Judge role	*"You are an expert US physician." Score every candidate independently using only the question or scenario and the blinded candidates; do not infer or identify which model produced a candidate.*
Ranking	*After scoring, rank all candidates from best- to worst-performing using overall clinical judgment; the ranking need not follow the score sums.*